

# **Project Report — Hyper-Personalized Customer Support AI**

## **1. Introduction**

Modern customers expect instant, context-aware support. Traditional chatbots fail because they:

- Give generic responses
  - Cannot personalize based on user history
  - Cannot understand live store data
  - Cannot integrate real-world context such as location

This project solves that gap by building a **Hyper-Personalized Customer Support AI** designed for retail and café environments.

The system provides:

- Personalized responses based on user profile
  - Live store information (inventory, offers, opening status)
  - Location-aware suggestions
  - Secure PII-masked LLM integration
  - Retrieval-Augmented Generation (RAG) for deeper personalization

The platform serves as an intelligent store assistant capable of elevating customer experience while reducing human support load.

## 2. Directory Structure

```

├── data/                      # Application datasets
│   ├── users.md               # User profiles (markdown)
│   ├── store.md                # Store metadata (markdown)
│   └── user_history/           # Conversation memory files

├── models/                    # Pretrained models (if any)

├── static/                    # Frontend UI (HTML/CSS/JS)
│   └── index.html              # Chat interface

└── venv/                      # Python virtual environment

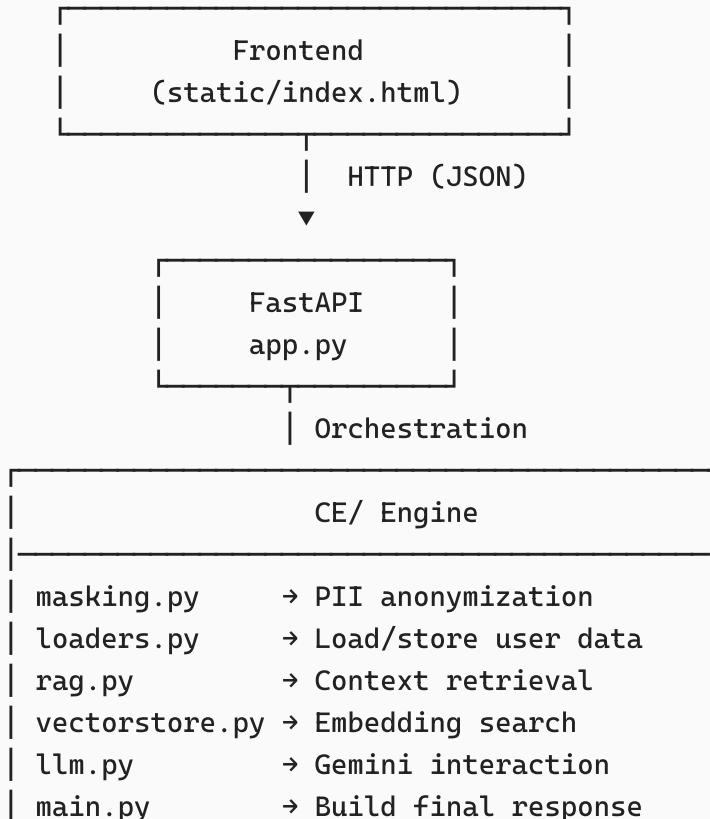
    ├── .env                   # Environment variables (Gemini API key)
    ├── .gitignore              # Ignore unneeded files
    ├── app.py                  # FastAPI server (main backend entry point)
    ├── README.md               # Project overview
    └── requirements.txt         # Dependencies

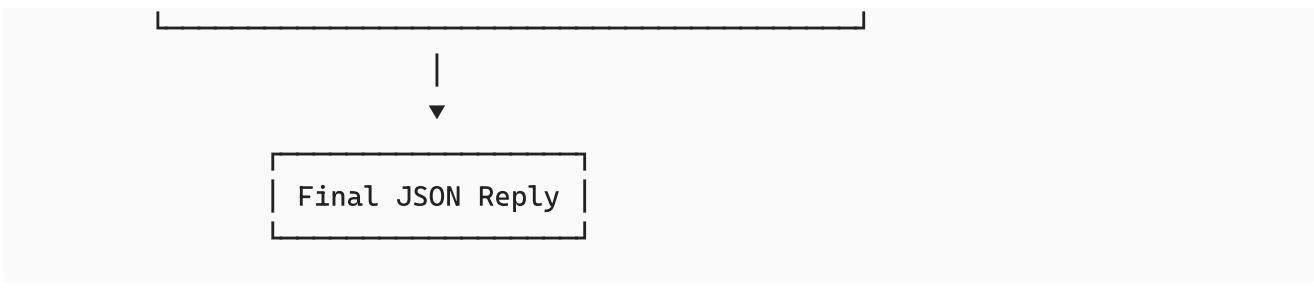
```

## 3. Core System Architecture

The system follows a **modular, production-friendly architecture**, similar to real microservices-based backend AI applications.

### 3.1 Architecture Diagram





## 4. Component Descriptions

### 4.1 CE/ — Core AI Engine

This folder contains all backend logic written in a clean, modular pattern.

#### **llm.py**

Centralized LLM module:

- Initializes Gemini API
- Builds prompts
- Fetches AI responses
- Handles errors & fallback messaging

#### **masking.py**

Responsible for masking:

- Phone numbers
- Email addresses
- Credit card numbers

Ensures **no raw PII** ever reaches the LLM.

#### **rag.py**

Handles RAG operations:

- PDF or markdown ingestion
- Chunk generation
- Embedding computation
- Context retrieval

#### **vectorstore.py**

Wrapper around FAISS for:

- Index creation
- KNN search
- Index persistence

## **loaders.py**

Reads and structures:

- User profiles
- Store details
- Inventory files

## **main.py**

This is the “brain” of the system:

- Merges user data + RAG + location + LLM
  - Formats the final response
  - Adds personalization into every message
- 

# **5. Data Layer**

## **users.md**

Markdown representation of user profiles (name, preferences, loyalty status).

## **store.md**

Markdown file describing the store location, offers, timings.

## **user\_history/**

Directory storing per-user conversation context to enable:

- Memory
  - Preference tracking
  - Better personalization
- 

# **6. Frontend UI**

The UI is built using:

- **TailwindCSS** (design system)

- **Vanilla JavaScript** (lightweight)
- **Smooth animations**
- **Starbucks-themed color palette**

Frontend features:

- Real chat UI
  - Quick actions (e.g., “I’m cold”)
  - Store info panel
  - Live response rendering
  - Typing indicator
- 

## 7. Backend (FastAPI)

**Features:**

- `/api/chat` → Main conversational endpoint
- `/api/health` → Service status
- PII masking
- Conversational memory
- Personalized replies
- Location-aware suggestions

**Technologies:**

- FastAPI
  - Gemini 2.5 Flash
  - Geopy
  - FAISS
  - Sentence Transformers
  - Dotenv
- 

## 8. Running the Application

**Step 1: Install dependencies**

```
pip install -r requirements.txt
```

**Step 2: Set your Gemini API key**

```
echo GEMINI_API_KEY=your_key_here > .env
```

## Step 3: Start the backend

```
uvicorn app.main:app --reload
```

## Step 4: Open the frontend

Open:

```
static/index.html
```

## 9. Key Technical Highlights

### Hyper-personalized responses

Based on:

- User identity
- Past chats
- Favorite drinks
- Loyalty tier
- Offers and inventory

### Location awareness

Bot uses geodesic distance to estimate how far the user is from the store.

### Real RAG pipeline

Uses chunking & embeddings to retrieve:

- Order history
- Preference notes
- Store policies

### Production-level privacy

PII masking ensures compliance with enterprise rules.

### Clean, modular architecture

Separates:

- LLM logic
- RAG logic
- Data loading
- Masking
- API layer

## Lightweight & fast

Optimized for hackathon-style rapid deployment.

---

## 10. Conclusion

This project demonstrates a **real-world, production-ready** solution to customer experience automation.

It integrates:

- AI
- RAG
- Real-time store context
- Geolocation
- Secure data handling
- A modern UI